

Notes vault — agent rules

This is a **learning / operator notes** vault (not a live attack tree).

HARD DEFAULT: every substantive deliverable follows the **ship pipeline** below.
Do not skip steps unless the operator explicitly opts out with the skip phrases.

Ship pipeline (mandatory — non-negotiable)

Whenever you **create, rewrite, expand, polish, or finish** work in this vault (study notes, guides, checklists, templates, operator docs, README teaching sections, eval harness files, or any other artifact the operator will keep):

1. CONTENT GEN draft / write the file(s) on disk
2. CONTENT EVAL run content_eval agent/skill 3× (see loop)
3. GIT commit + push to origin (never force-add logs/)
4. MAIL email markdown + PDF (when prose) to study inbox

A task is **not done** until steps that apply have finished (or the operator used an explicit skip phrase). Saying “here is the draft” in chat without git + mail is incomplete.

Step 1 — Content gen

- Write real files under the vault (prefer AI/ , topic folders, etc.).
- For **study / course topics**, complete the **study research pack** first (official docs + gold-mine blogs), then draft.

Step 2 — Content gen ↔ content_eval 3× agent loop

After the draft exists on disk:

1. Load **content_eval** (skill /content-eval or agent content_eval).
2. Run **exactly three sequential passes** on the content:

Pass	Lens	Job
1	slop_chop	Kill AI fluff; raise density
2	first_principles	Definition, mechanism, high-level map, beginner-clear

Pass	Lens	Job
3	core_questions	3-7 core Qs answered or marked gap; force clarity

1. On **FAIL** for a pass: rewrite (max 2 rewrites per pass) → re-run that lens.
2. Overall must be **SHIP** (or **BLOCKED** with ≤ 5 human questions).
3. Report the three verdicts + overall before claiming done.

Sources:

piece	path
agent	.grok/agents/content_eval.md (subagent_type: content_eval)
skill	.grok/skills/content-eval/SKILL.md → /content-eval
rule	.grok/rules/content-eval-mandatory.md

Skip content_eval only if: operator says `skip content-eval / raw dump ok / no eval`, or pure mechanical change (typo/path/git metadata) with no teaching prose.

Step 3 — Git (always)

After content is SHIP (or non-prose change is ready):

1. `git status / git diff / recent log` style as usual.
2. Stage **only** intended artifacts (never `git add -f logs/` or secrets).
3. Commit with a clear message (why this change).
4. **Push** to `origin` (SSH/HTTPS as configured). `Rebase/pull` if needed; resolve conflicts honestly.
5. Report commit hash + push result to the operator.

Skip git only if: operator says `skip git / no push / local only`.

If author/identity or auth fails: fix or report the exact blocker — do not silently leave work uncommitted when git was required.

Step 4 — Mail (always for keepable artifacts)

Default: **email every keepable deliverable** to the study inbox via pwnjournal SMTP.

```
# typical (paths adjust to workspace)
python3 /home/kali/HTB/PwnJournal/scripts/md_to_pdf.py PATH.md -o PATH.pdf
python3 /home/kali/HTB/PwnJournal/scripts/send_report_email.py \
  --subject "[Notes] <short title>" \
  --body "<one-line what shipped>" \
  PATH.pdf PATH.md
```

- Config: `~/.config/pwnjournal/smtp.env` (or `PWNJOURNAL_SMTP_ENV / .env.smtp`).
- Default To: `SMTP_TO` (study inbox, e.g. `dailyupdatesforstudies@gmail.com`).

- Attach **PDF + markdown** when the deliverable is prose notes.
- For multi-file packs, attach the full set in one mail when practical.
- If mail fails: still keep files + git; report the error (do not pretend mailed).

Skip mail only if: operator says skip mail / no email / don't mail.

Study research pack (mandatory on every study topic)

Whenever you **create or expand study notes** for a course module, product topic, or learning unit (Associate modules, platform foundations, etc.):

1. Fetch official product docs (Claude / Anthropic when the course is Claude)

Use web search + open primary pages. Prefer:

- support.claude.com Help Center
- [platform.claude.com](https://platform.claude.com/docs) / docs.anthropic.com platform docs
- anthropic.com/engineering and anthropic.com/news for mechanism-level posts
- Plan/UI limits: re-fetch (they change); never invent caps

Capture **links + one-line so-what** under **Official Claude / Anthropic references** (or vendor-equivalent if the topic is not Claude).

2. Fetch gold-mine sources outside product docs

Reputable **first-principles** material that is *not* a click-path UI tutorial:

Prefer	Examples
Vendor engineering (mechanism)	Anthropic / OpenAI engineering
Known practitioners	Simon Willison, Addy Osmani, similar long-form
Concept explainers	Context eng vs prompt eng, agent memory, tools
Avoid as sole authority	SEO listicles, anonymous “top 10”, undated affiliate posts

Every study note must include:

```
## Gold-mine blogs (outside Claude product docs)
```

(or “outside <vendor> product docs”). Each row: link + why it is gold.

3. Do not restrict examples to one vendor

Map concepts to ChatGPT / Gemini / Grok / local agents when principles transfer.

4. Then run the full ship pipeline

Research → draft → content_eval 3× → git push → mail.

Skip research pack only if: operator says skip research / course text only.

Rule file: .grok/rules/study-sources-mandatory.md.

Skip phrases (only valid opt-outs)

Phrase	Skips
skip content-eval / raw dump ok / no eval	Step 2
skip git / no push / local only	Step 3
skip mail / no email / don't mail	Step 4
skip research / course text only	Study research pack
ship pipeline off	All of steps 2-4 for that turn only (still write files if asked)

No skip phrase → **run the full pipeline.**

Tone for this vault

- frugal, copy-paste useful, just-in-time why
- first principles over tour-guide fluff
- authorized-testing scope only when discussing attacks